

TP 1

Adan Rodriguez

DataSet usado: breast cancer diagnosis

Objetivo: saber si a partir de ciertos detalles del tumor, intentar predecir si un tumor podra ser maligno o benigno

```
In [257... import pandas as pd
import statsmodels.api as sm
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import metrics
```

```
In [258... df = pd.read_csv('data.csv')
df.head()
```

```
Out[258]:
```

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness
0	842302	M	17.99	10.38	122.80	1001.0	0.1184
1	842517	M	20.57	17.77	132.90	1326.0	0.1425
2	84300903	M	19.69	21.25	130.00	1203.0	0.1875
3	84348301	M	11.42	20.38	77.58	386.1	0.2763
4	84358402	M	20.29	14.34	135.10	1297.0	0.1634

5 rows × 33 columns

Ahora limpiare el dataset de modo que omita campos como uri, artists_names etc y se persistira solo datos cuantitativos

```
In [259... df = df.drop(['id'], axis=1)
df.head()
```

```
Out[259]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness
0	M	17.99	10.38	122.80	1001.0	0.1184
1	M	20.57	17.77	132.90	1326.0	0.1425
2	M	19.69	21.25	130.00	1203.0	0.1875
3	M	11.42	20.38	77.58	386.1	0.2763
4	M	20.29	14.34	135.10	1297.0	0.1634

5 rows × 32 columns

ya que la ultima columna esta constituida por NaN, la eliminamos del dataset

```
In [260]: df = df.dropna(axis=1)
df.head()
```

```
Out[260]:
```

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_
0	M	17.99	10.38	122.80	1001.0	0
1	M	20.57	17.77	132.90	1326.0	0.
2	M	19.69	21.25	130.00	1203.0	0.
3	M	11.42	20.38	77.58	386.1	0.
4	M	20.29	14.34	135.10	1297.0	0.

5 rows × 31 columns

```
In [261]: '''Vemos si hay datos nulos en nuestro dataset'''
df.isnull().any()
```

```
Out[261]:
```

diagnosis	False
radius_mean	False
texture_mean	False
perimeter_mean	False
area_mean	False
smoothness_mean	False
compactness_mean	False
concavity_mean	False
concave points_mean	False
symmetry_mean	False
fractal_dimension_mean	False
radius_se	False
texture_se	False
perimeter_se	False
area_se	False
smoothness_se	False
compactness_se	False
concavity_se	False
concave points_se	False
symmetry_se	False
fractal_dimension_se	False
radius_worst	False
texture_worst	False
perimeter_worst	False
area_worst	False
smoothness_worst	False
compactness_worst	False
concavity_worst	False
concave points_worst	False
symmetry_worst	False
fractal_dimension_worst	False
dtype:	bool

```
In [262]: #Resumen del dataframe
df.describe()
```

Out[262]:

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
count	569.000000	569.000000	569.000000	569.000000	569.000000
mean	14.127292	19.289649	91.969033	654.889104	0.096360
std	3.524049	4.301036	24.298981	351.914129	0.014064
min	6.981000	9.710000	43.790000	143.500000	0.052630
25%	11.700000	16.170000	75.170000	420.300000	0.086370
50%	13.370000	18.840000	86.240000	551.100000	0.095870
75%	15.780000	21.800000	104.100000	782.700000	0.105300
max	28.110000	39.280000	188.500000	2501.000000	0.163400

8 rows × 30 columns

Ahora traduciremos el valor categorico M o B a uno numerico 1 o 0

```
In [263... from sklearn.preprocessing import LabelEncoder, OneHotEncoder
labelencoder_X = LabelEncoder()

#Transformo la columna
df.iloc[:, 0] = labelencoder_X.fit_transform(df.iloc[:, 0])

df.head()
```

Out[263]:

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_
0	1	17.99	10.38	122.80	1001.0	0
1	1	20.57	17.77	132.90	1326.0	0
2	1	19.69	21.25	130.00	1203.0	0
3	1	11.42	20.38	77.58	386.1	0
4	1	20.29	14.34	135.10	1297.0	0

5 rows × 31 columns

```
In [264... #Separamos la variable independiente
y = pd.DataFrame(df.iloc[:, 0])
X = df.drop('diagnosis', axis=1)

X_df = X
```

Ahora para ver el p valor de cada campo le agregamos una fila de 1 a la matriz

```
In [265... # agregamos la columna de 1 al conjunto X original
X = np.append(arr = np.ones((569,1)).astype(int), values = X, axis = 1)
```

```
In [266... def EliminacionBackward(x, sl):
    numVars = len(x[0])
```

```

for i in range(0, numVars):
    regressor_OLS = sm.OLS(y, x.tolist()).fit()
    maxVar = max(regressor_OLS.pvalues)
    if maxVar > sl:
        for j in range(0, numVars - i):
            if (regressor_OLS.pvalues[j] == maxVar):
                x = np.delete(x, j, 1)
    regressor_OLS.summary()
return x

```

In [267... *#Eliminamos las variables con mayor valor p automaticamente y mostramos e*

```

SL = 0.05
y = y.astype(int)
X_opt = X

X_Modelado = EliminacionBackward(X_opt, SL)

pd.DataFrame(X_Modelado).head()

```

Out[267]:

	0	1	2	3	4	5	6	7	8	9
0	1.0	17.99	0.27760	0.14710	1.0950	0.006399	0.05373	0.01587	25.38	17.33
1	1.0	20.57	0.07864	0.07017	0.5435	0.005225	0.01860	0.01340	24.99	23.41
2	1.0	19.69	0.15990	0.12790	0.7456	0.006150	0.03832	0.02058	23.57	25.53
3	1.0	11.42	0.28390	0.10520	0.4956	0.009110	0.05661	0.01867	14.91	26.50
4	1.0	20.29	0.13280	0.10430	0.7572	0.011490	0.05688	0.01885	22.54	16.67

In [268... *#Eliminacion hecha a mano para comparar con la eliminacion automatica*

```

X_opt = X[:, [0, 1, 6, 8, 11, 15, 17, 18, 21, 22, 24, 27, 29, 30]]
#Con este X auxiliar veremos y eliminaremos las columnas que menos influy
# X_opt, al inicio, tomará todas las filas y cada una de las columnas del

#convierto columna a tipo numerico
y = y.astype(int)

# tecnica OLS
SL = 0.05
regression_OLS = sm.OLS(endog = y, exog = X_opt.tolist()).fit()
# observar el p valor en el sumario
regression_OLS.summary()

```

Out[268]:

OLS Regression Results

Dep. Variable:	diagnosis			R-squared:	0.771	
Model:	OLS			Adj. R-squared:	0.766	
Method:	Least Squares			F-statistic:	144.1	
Date:	Sat, 27 May 2023			Prob (F-statistic):	3.34e-168	
Time:	14:34:02			Log-Likelihood:	26.008	
No. Observations:	569			AIC:	-24.02	
Df Residuals:	555			BIC:	36.80	
Df Model:	13					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
const	-2.0783	0.182	-11.449	0.000	-2.435	-1.722
x1	-0.0340	0.015	-2.224	0.027	-0.064	-0.004
x2	-3.3622	0.552	-6.092	0.000	-4.446	-2.278
x3	4.9320	0.943	5.230	0.000	3.080	6.784
x4	0.2147	0.069	3.105	0.002	0.079	0.351
x5	19.6509	4.335	4.533	0.000	11.135	28.167
x6	-3.0397	0.747	-4.067	0.000	-4.508	-1.571
x7	7.7568	3.348	2.317	0.021	1.180	14.334
x8	0.1657	0.021	8.041	0.000	0.125	0.206
x9	0.0102	0.002	5.552	0.000	0.007	0.014
x10	-0.0009	0.000	-7.663	0.000	-0.001	-0.001
x11	0.6533	0.144	4.537	0.000	0.370	0.936
x12	0.7751	0.213	3.634	0.000	0.356	1.194
x13	3.3066	1.139	2.903	0.004	1.070	5.544
Omnibus:	27.517	Durbin-Watson:		1.777		
Prob(Omnibus):	0.000	Jarque-Bera (JB):		30.209		
Skew:	0.554	Prob(JB):		2.76e-07		
Kurtosis:	3.218	Cond. No.		4.71e+05		

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 4.71e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [269... #Ahora asignamos el dataframe real con los campos que sabemos que su valor no es 0
X_df = X_df.iloc[:, [0, 5, 7, 10, 14, 16, 17, 20, 21, 23, 26, 28, 29]]
```

X_df

Out[269]:


	radius_mean	compactness_mean	concave points_mean	radius_se	smoothness_se	co
0	17.99	0.27760	0.14710	1.0950	0.006399	
1	20.57	0.07864	0.07017	0.5435	0.005225	
2	19.69	0.15990	0.12790	0.7456	0.006150	
3	11.42	0.28390	0.10520	0.4956	0.009110	
4	20.29	0.13280	0.10430	0.7572	0.011490	
...
564	21.56	0.11590	0.13890	1.1760	0.010300	
565	20.13	0.10340	0.09791	0.7655	0.005769	
566	16.60	0.10230	0.05302	0.4564	0.005903	
567	20.60	0.27700	0.15200	0.7260	0.006522	
568	7.76	0.04362	0.00000	0.3857	0.007189	

569 rows × 13 columns

In [270]...

```
#Se muestra la matriz de correlacion resultante
X_aux = X_df.iloc[:, :]
X_aux.insert(0, 'diagnosis', y)
corr = pd.DataFrame(X_aux).corr()
plt.subplots(figsize=(12,8))
sns.heatmap(corr, xticklabels=corr.columns, yticklabels=corr.columns, ann
            cmap=sns.diverging_palette(240, 10, as_cmap=True))
```

Out[270]: <Axes: >



```
In [346]: X_aux.describe()
```

	diagnosis	radius_mean	compactness_mean	concave points_mean	radius_se	smc
count	569.000000	569.000000	569.000000	569.000000	569.000000	
mean	0.372583	14.127292	0.104341	0.048919	0.405172	
std	0.483918	3.524049	0.052813	0.038803	0.277313	
min	0.000000	6.981000	0.019380	0.000000	0.111500	
25%	0.000000	11.700000	0.064920	0.020310	0.232400	
50%	0.000000	13.370000	0.092630	0.033500	0.324200	
75%	1.000000	15.780000	0.130400	0.074000	0.478900	
max	1.000000	28.110000	0.345400	0.201200	2.873000	

- Definimos el significado de los campos:
- diagnosis: diagnostico final del tumor (M/1: Maligno, B/0: Benigno)
 - radius_mean: media de las distancias desde el medio a varios puntos del tumor
 - compactness_mean: el perimetro cuadrado dividido por el area-1
 - concave points_mean: media del numero de porciones concavas del contorno
 - radius_se: error estandar de la medida del radio
 - smoothness_se: variacion local en longitud de radio
 - concavity_se: severidad de las porciones concavas del contorno
 - concave points_se: error estandar de las porciones concavas

radius_worst: distancia mas amplia entre el centro y algun punto del tumor

texture_worst: peor desviacion de los valores en escala de grises

area_worst: mayor area del tumor

concavity_worst: concavidad mas severa del contorno

symmetry_worst: peor simetria del tumor

fractal_dimension_worst: peor irregularidad de los contornos del tumor

```
In [271]: corr['diagnosis'].abs().sort_values(ascending=False)
```

```
Out[271]: diagnosis                1.000000
concave points_mean             0.776614
radius_worst                    0.776454
area_worst                      0.733825
radius_mean                     0.730029
concavity_worst                 0.659610
compactness_mean                0.596534
radius_se                       0.567134
texture_worst                   0.456903
symmetry_worst                  0.416294
concave points_se               0.408042
fractal_dimension_worst         0.323872
concavity_se                    0.253730
smoothness_se                   0.067016
Name: diagnosis, dtype: float64
```

Vemos la correlacion entre el diagnostico y las variables

Ahora reacomodamos nuestra matriz:

```
In [272]: #Asigno la matriz resultante a la matriz original de variables dependientes
X = X_df
X
```

```
Out[272]:
```

	radius_mean	compactness_mean	concave points_mean	radius_se	smoothness_se	co
0	17.99	0.27760	0.14710	1.0950	0.006399	
1	20.57	0.07864	0.07017	0.5435	0.005225	
2	19.69	0.15990	0.12790	0.7456	0.006150	
3	11.42	0.28390	0.10520	0.4956	0.009110	
4	20.29	0.13280	0.10430	0.7572	0.011490	
...
564	21.56	0.11590	0.13890	1.1760	0.010300	
565	20.13	0.10340	0.09791	0.7655	0.005769	
566	16.60	0.10230	0.05302	0.4564	0.005903	
567	20.60	0.27700	0.15200	0.7260	0.006522	
568	7.76	0.04362	0.00000	0.3857	0.007189	

569 rows × 13 columns




```
In [273... #Separamos nuestros conjuntos
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
```

```
In [274... regressor = LinearRegression()
regressor.fit(X_train, y_train)
```

```
Out[274]: ▼ LinearRegression
LinearRegression()
```

```
In [275... y_pred = regressor.predict(X_test)

#Redondeamos los resultados para tener una respuesta binaria (tambien por
y_test = np.round(y_test, decimals=0, out=None)
y_pred = np.round(y_pred, decimals=0, out=None)
```

```
In [328... #Redimensionamos el array
y_test = np.array(y_test).reshape(-1)
y_pred = np.array(y_pred).reshape(-1)

#Cambiamos el tipo de valor a entero
y_pred = y_pred.astype(int)
y_test = y_test.astype(int)

#Creamos el dataframe para mostrar nuestra prediccion
df = pd.DataFrame({'Actual': y_test, 'Predicted': y_pred})

#Reemplazamos los valores numericos por categoricos para mejor visualizac
df['Actual'].replace(to_replace=0, value='B', regex=True, inplace=True)
df['Actual'].replace(to_replace=1, value='M', regex=True, inplace=True)

df['Predicted'].replace(to_replace=0, value='B', regex=True, inplace=True)
df['Predicted'].replace(to_replace=1, value='M', regex=True, inplace=True)
df['Predicted'].replace(to_replace=2, value='M', regex=True, inplace=True)

df_toShow = df.head(25)
df_toShow.head()
```

```
Out[328]:
```

	Actual	Predicted
0	M	M
1	B	B
2	B	B
3	B	B
4	B	B

```
In [330... def entropia(y):
    if isinstance(y, pd.Series):
        a = y.value_counts() / y.shape[0]
        entropy = np.sum(-a * np.log2(a + 1e-9))
        return entropy
    else:
        raise('el obj debe ser una serie Pandas')

entropia(X_aux.diagnosis)
```

Out[330]: 0.9526351195164697

Segun el calculo de la entropia, al estar cerca del 1 podemos afirmar que los datos varian bastante entre M y B

Ahora, veremos el conteo de diagnosticos con los que queriamos comparar nuestra prediccion

```
In [331...] graf1 = pd.crosstab(index = ['Diagnosis'], columns=df['Actual'], margins=
graf1
```

```
Out[331]:
```

	Actual	B	M	All
row_0				
Diagnosis	67	47	114	
All	67	47	114	

Conteo de los resultados de la prediccion

```
In [332...] graf2 = pd.crosstab(index = ['Diagnosis'], columns=df['Predicted'], margi
graf2
```


```
Out[332]:
```

	Predicted	B	M	All
row_0				
Diagnosis	70	44	114	
All	70	44	114	

Visualizamos los datos


```
In [338...] graf1[0:1].plot(kind='bar')
```

```
Out[338]: <Axes: xlabel='row_0'>
```



```
In [342...] graf2[0:1].plot(kind='bar')
```

```
Out[342]: <Axes: xlabel='row_0'>
```



Por ultimo calculamos el indice de error del algoritmo

```
In [343... print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, y_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, y_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

Mean Absolute Error: 0.043859649122807015

Mean Squared Error: 0.043859649122807015

Root Mean Squared Error: 0.20942695414584775

Viendo que nuestras metricas son cercanas al cero, podemos tener confianza de que nuestro algoritmo tendrá buenas predicciones

En conclusion, pudimos ver que en el diagnostico cancerigeno de un tumor, antes de hacer una biopsia para saber con certeza la naturaleza del mismo, la concavidad, el radio y el area influye a la hora de saber si es maligno o benigno, por lo que con un algoritmo como el propuesto, podriamos monitorear el cambio de estos valores en algun paciente para saber de antemano que resultados podria esperar, y asi estudiarlo mas fondo en pro de diagnosticarle lo mas pronto posible y tratarle en una etapa temprana